


Jan 22, 2026 Codex

# Testing Agent Skills Systematically with Evals

A practical guide to turning agent skills into something you can test, score, and improve over time.

Authors: Dominik Kundel, Gabriel Chua



When you're iterating on a skill for an agent like Codex, it's hard to tell whether you're actually improving it or just changing its behavior. One version feels faster, another seems more reliable, and then a regression slips in: the skill doesn't trigger, it skips a required step, or it leaves extra files behind.

At its core, a skill is an organized collection of prompts and instructions for an LLM. The most reliable way to improve a skill over time is to evaluate it the same way you would any other prompt for LLM applications.

*Evals* (short for *evaluations*) check whether a model's output, and the steps it took to produce it, match what you intended. Instead of asking "does this feel better?" (or relying on vibes), evals let you ask concrete questions like:

- Did the agent invoke the skill?

Ask AI